



Trinity College Dublin

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

01 – Memory and Addressing Modes

CS1022 – Introduction to Computing II

Dr Jonathan Dukes / jdukes@scss.tcd.ie
School of Computer Science and Statistics

LDR and STR

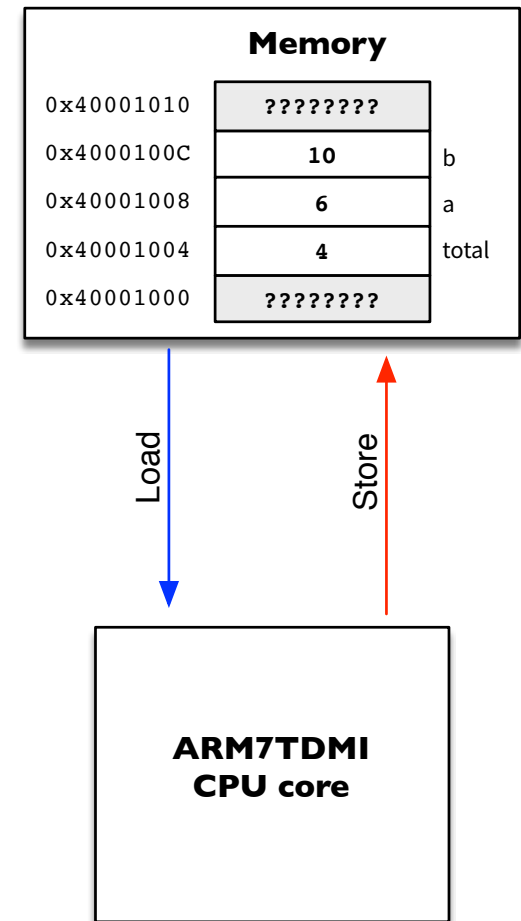
LDR and STR Reviewed

3

How many memory accesses are required to compute

$$\mathbf{total} = \mathbf{total} + (\mathbf{a} \times \mathbf{b})$$

where **total**, **a** and **b** are stored in memory at the addresses contained in **R0**, **R1** and **R2** respectively?



LDR and STR Reviewed

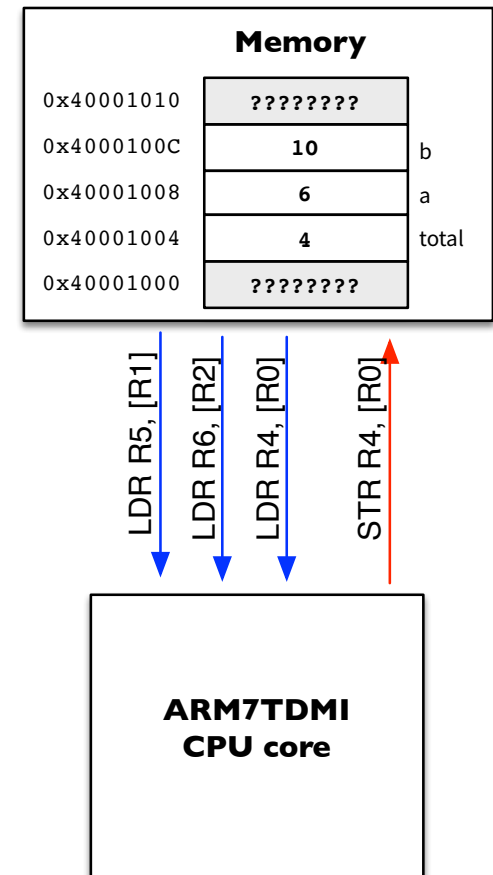
4

How many memory accesses are required to compute **total = total + (a × b)**, where **total**, **a** and **b** are stored in memory at the addresses contained in **R0**, **R1** and **R2** respectively?

```
LDR    R5, [R1]      ; Load a
LDR    R6, [R2]      ; Load b
MUL    R5, R6, R5     ; tmp = a * b

LDR    R4, [R0]      ; Load total
ADD    R4, R4, R5     ; total = total + (tmp)

STR    R4, [R0]      ; Store total back to memory
```



Assembler Directives for Memory Layout

5

AREA

Declares a code or data section and defines its attributes (CODE, DATA, READONLY, READWRITE, etc.)

Section – an indivisible sequence of instructions and/or region of memory used to store data that is arranged in memory by the linker

SPACE

Tells the assembler to set aside space in memory to contain data

Does not initialise anything

(We would have to write code to initialise data that we want to modify later)

DCB, DCW, DCD

Initialises ROM with constant Byte-, Halfword- or Word-size values

```
AREA    globals, DATA, READWRITE
L0 SPACE    4                      ; will be placed at 0x400000000

AREA    RESET, CODE, READONLY
ENTRY

LDR     R0, =L0
LDR     R1, =L1
LDR     R2, =L2
LDR     R4, [R1]
LDR     R5, [R2]
ADD     R4, R4, R5
STR     R4, [R0]


STOP
B       STOP

L1 DCD     6                      ; will be placed after our code
L2 DCD     7

END
```

Load / Store Bytes, Half-words and Words

6



Byte, half-word and word at address 0xA00000D8

```
LDR r0, =0x400000D8
LDRB r1, [r0]
```

00	00	00	32
----	----	----	----

```
LDR r0, =0x400000D8
LDRH r1, [r0]
```

00	00	04	32
----	----	----	----

```
LDR r0, =0x400000D8
LDR r1, [r0]
```

EB	85	04	32
----	----	----	----

Alignment

7

ARM7TDMI expects all memory accesses to be **aligned**

Examples:

Address Examples

Word aligned	0x00000000, 0x00001008, 0x4100000C
Not word aligned	0x00000001, 0x00001006, 0x4100000F
Half-word aligned	0x00000000, 0x00001002, 0x4100000A
Not half-word aligned	0x00000003, 0x00001001, 0x4100000B

See ARM Architecture Reference Manual Section A2.8

Unaligned accesses are permitted but the result is unlikely to be what was intended

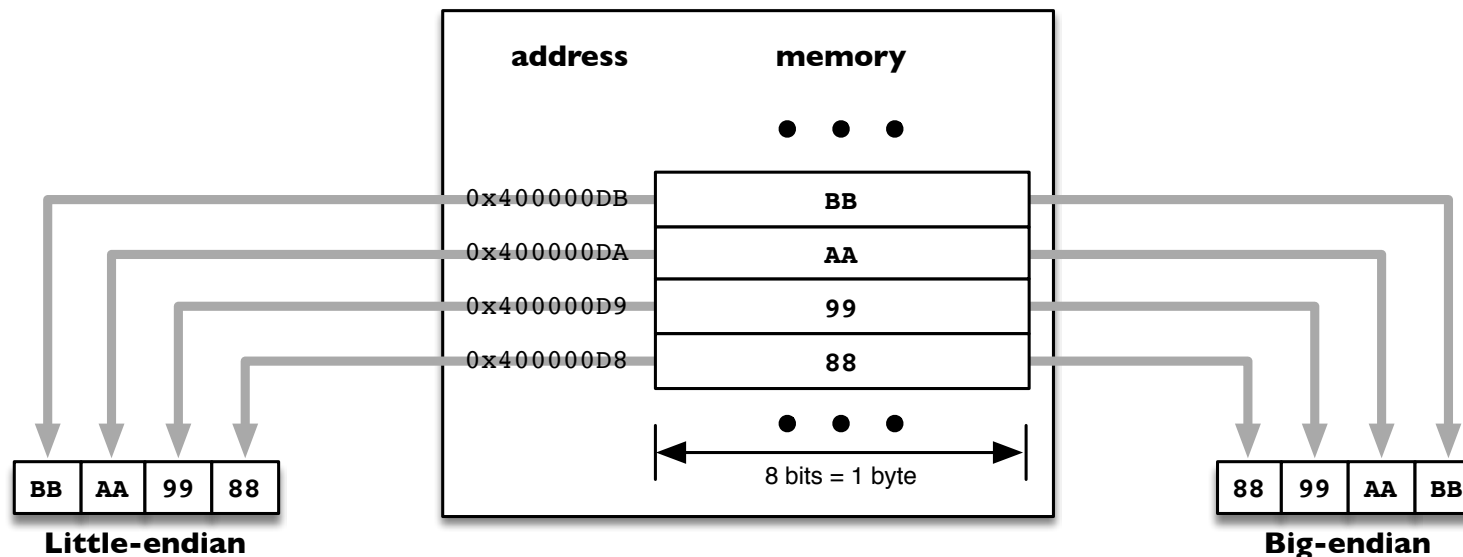
Unaligned accesses are supported by later ARM architecture versions

Endian-ness

8

Little-endian byte ordering – least-significant byte of word or half-word stored at lower address in memory

Big-endian byte ordering – most-significant byte of word or half-word stored at lower address in memory



Loading and storing signed values

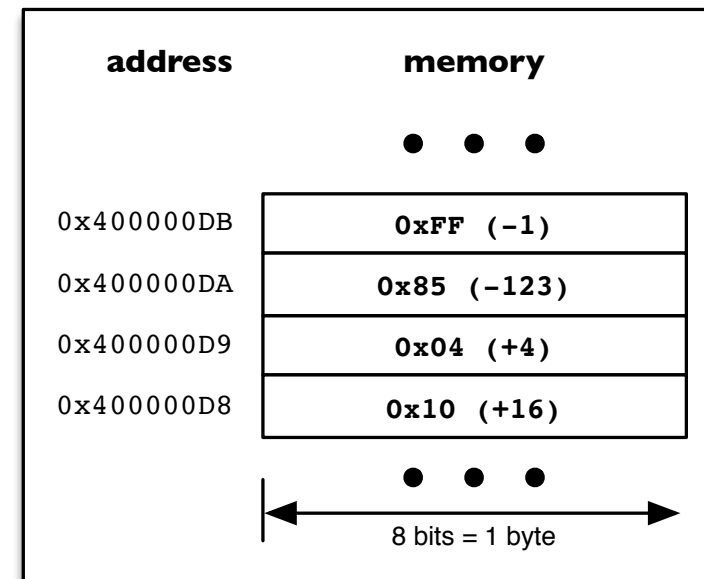
9

Consider the four byte-sized signed 2's complement values stored in memory on the right

After executing the pair of instructions below, what is the correct signed decimal interpretation of the value loaded in R1?

```
LDR    r0, =0x400000D8
LDRB   r1, [r0]
```

- (a) -16
- (b) +240
- (c) +16
- (d) -240



Loading and storing signed values

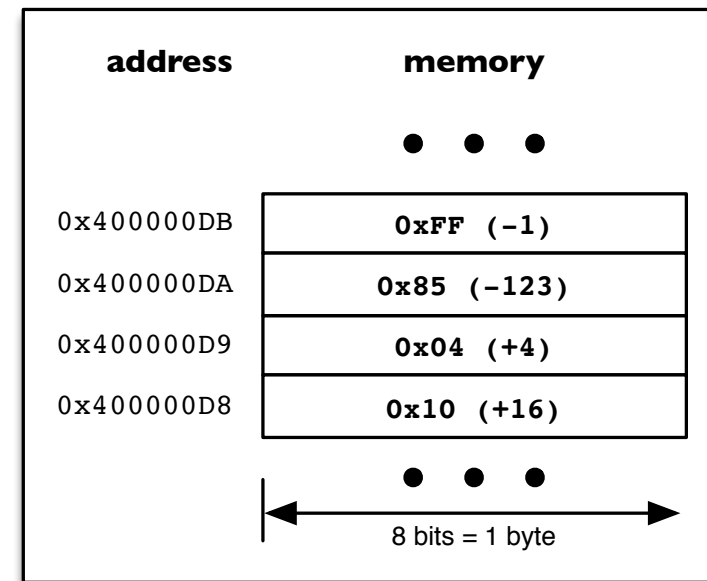
10

Consider the four byte-sized signed 2's complement values stored in memory on the right

After executing the pair of instructions below, what is the correct signed decimal interpretation of the value loaded in R1?

```
LDR    r0, =0x400000DB
LDRB   r1, [r0]
```

- (a) +1
- (b) +255
- (c) -1
- (d) -255




Loading Signed Bytes and Half-words

11

Sign extension performed when loading signed bytes or half-words to facilitate correct subsequent 32-bit signed arithmetic

```
LDR    r0, =0x400000D8
LDRSB  r1, [r0]
```



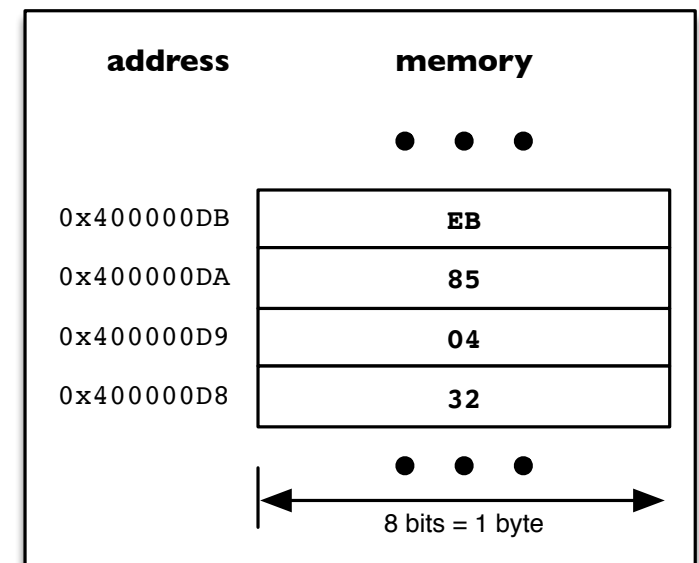
```
LDR    r0, =0x400000D8
LDRSH  r1, [r0]
```



```
LDR    r0, =0x400000DA
LDRSB  r1, [r0]
```



```
LDR    r0, =0x400000DA
LDRSH  r1, [r0]
```



Loading and storing signed values


12

Consider the four byte-sized signed 2's complement values stored in memory on the right

After executing the pair of instructions below, what is the correct signed decimal interpretation of the value loaded in R1?

```
LDR    r0, =0x400000DA
LDRB   r1, [r0]
```

- (a) +123
- (b) -123
- (c) +133
- (d) -252



Loading and storing signed values


13

Consider the four byte-sized signed 2's complement values stored in memory on the right

After executing the pair of instructions below, what is the correct signed decimal interpretation of the value loaded in R1?

```
LDR    r0, =0x400000DA
LDRSB  r1, [r0]
```

- (a) +123
- (b) -123
- (c) +133
- (d) -252



Loading and storing signed values

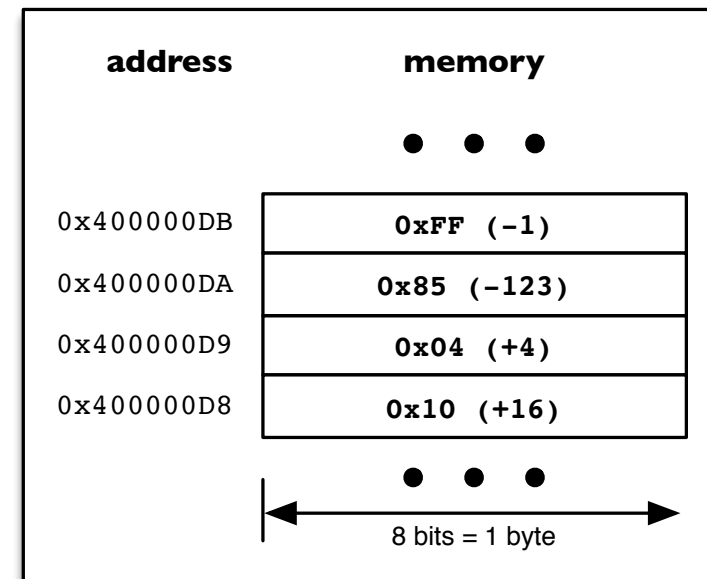
14

Consider the four byte-sized signed 2's complement values stored in memory on the right

After executing the pair of instructions below, what is the correct signed decimal interpretation of the value loaded in R1?

```
LDR    r0, =0x400000D9
LDRSB  r1, [r0]
```


- (a) -4
- (b) -252
- (c) +252
- (d) +4



Addressing Modes

Addressing Modes

16



The syntax **[R1]** is just one of many ways that we can specify the the memory address that we want to access using LDR or STR

Remember: **[R1]** tells the processor to access the value in memory at the address contained in register R1

The syntax **[R1]** corresponds to an **Addressing Mode**

[R1] is an abbreviated form of **[R1, #0]** (the #0 is implied if omitted)

The address accessed by LDR or STR is called the **Effective Address (EA)**

Immediate Offset Addressing

17

Addressing Mode Syntax	Operation	Example
[Rn, #offset]	EA = Rn + offset	LDR R0, [R1, #4]
[Rn]	EA = Rn + #0 (<i>#0 is assumed</i>)	LDR R0, [R1]

Effective Address is calculated by adding **offset** to the address in the base register **Rn** (note: offset may be negative)

Base register **Rn** is not changed

Example: load three consecutive word-size values from memory into registers R4, R5 and R6, beginning at the address contained in R0

```
LDR    R0, =label    ; Initialise base register
LDR    R4, [R0]       ; R4 = Memory.word[R0 + 0] (default = 0)
LDR    R5, [R0, #4]   ; R5 = Memory.word[R0 + 4]
LDR    R6, [R0, #8]   ; R6 = Memory.word[R0 + 8]
```

Register Offset Addressing

18

Addressing Mode Syntax	Operation	Example
[Rn, Rm]	$EA = Rn + Rm$	LDR R0, [R1, R2]

Effective Address is calculated by adding offset in **Rm** to the address in the base register **Rn**

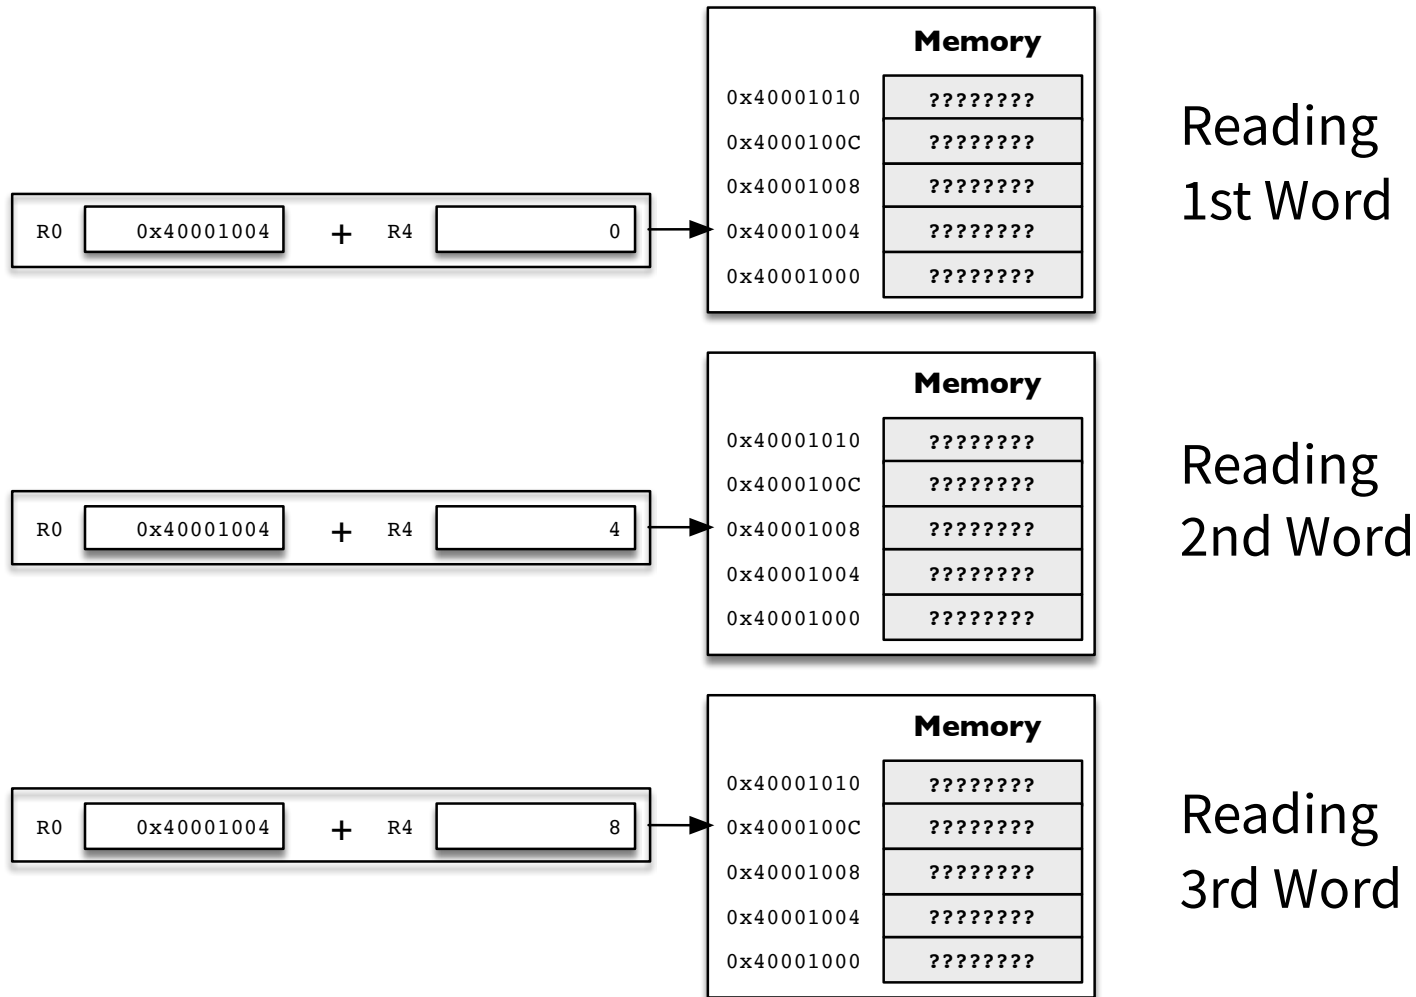
Base register **Rn** and offset register **Rm** are not changed

Example: load three consecutive word values from memory into registers R1, R2 and R3 beginning at the address in R0

```
LDR    r0, =label      ; Initialise base register
LDR    r4, =0           ; Initialise offset register = 0
LDR    r1, [r0, r4]     ; r1 = Memory.word[r0 + r4]
ADD    r4, r4, #4       ; r4 = r4 + 4
LDR    r2, [r0, r4]     ; r2 = Memory.word[r0 + r4]
ADD    r4, r4, #4       ; r4 = r4 + 4
LDR    r3, [r0, r4]     ; r3 = Memory.word[r0 + r4]
```

Register Offset Addressing

19



Example: Convert to UPPER CASE (using Register Offset Addressing)

20

```
LDR    r1, =TESTSTR    ; address = TESTSTR
LDR    r2, =0           ; index = 0

LDRB   r0, [r1, r2]     ; char = Memory.Byte[address + index]
L1
CMP     r0, #0           ; while ( char != 0 )
BEQ     L3               ; {
CMP     r0, #'a'         ; if (char >= 'a'
BLO     L2               ;   &&
CMP     r0, #'z'         ;   char <= 'z' )
BHI     L2               ; {
BIC     r0, #0x00000020 ;   char = char AND NOT 0x00000020
STRB    r0, [r1, r2]     ;   Memory.byte[address + index] = char
L2
        ; }
ADD     r2, r2, #1       ; index = index + 1
LDRB    r0, [r1, r2]     ; char = Memory.Byte[address + index]
B       L1               ; }
L3
        ;
```

Example: Convert to UPPER CASE (using Register Offset Addressing)

21

```
LDR    r1, =TESTSTR    ; address = TESTSTR
LDR    r2, =0           ; index = 0

L1
  LDRB  r0, [r1, r2]     ; while ( (char = Memory.Byte[address + index])
  CMP   r0, #0           ;             != 0 )
  BEQ   L3               ; {
  CMP   r0, #'a'         ;   if (char >= 'a'
  BLO   L2               ;   &&
  CMP   r0, #'z'         ;   char <= 'z')
  BHI   L2               ; {
  BIC   r0, #0x00000020 ;   char = char AND NOT 0x00000020
  STRB  r0, [r1, r2]     ;   Memory.byte[address + index] = char
L2
  ADD   r2, r2, #1       ;   index = index + 1
LDRB  r0, [r1, r2]     ; char = Memory.Byte[address + index]
  B     L1               ; }
L3
```

By moving the label *whStr* to include the top LDRB, we can eliminate the bottom LDRB. We will use this construct from now on.

pseudo-code in blue evaluates to the newly loaded value of char (i.e. just like Java, C, etc..!)

no longer needed!

LDR and STR Addressing Modes

Mode	Pseudo-code Operation *	Example
Offset modes – the base address register <i>is not</i> modified		
[Rn]	EA = Rn	LDR R0, [R1]
[Rn, #offset]	EA = Rn + offset	LDR R0, [R1, #4]
[Rn, Rm]	EA = Rn + Rm	LDR R0, [R1, R2]
[Rn, Rm, LSL #shift]	EA = Rn + (Rm << shift)	LDR R0, [R1, R2, LSL #2]
[Rn, Rm, LSR #shift]	EA = Rn + (Rm >> shift)	LDR R0, [R1, R2, LSR #2]
Pre-indexed modes – the base address register <i>is</i> modified <i>before</i> accessing memory		
[Rn, #offset]!	Rn = Rn + offset EA = Rn	LDR R0, [R1, #4]!
[Rn, Rm]!	Rn = Rn + Rm EA = Rn	LDR R0, [R1, R2]!
[Rn, Rm, LSL #shift]!	Rn = Rn + (Rm << shift) EA = Rn	LDR R0, [R1, R2, LSL #2]!
[Rn, Rm, LSR #shift]!	Rn = Rn + (Rm >> shift) EA = Rn	LDR R0, [R1, R2, LSR #2]!
Post-indexed modes – the base address register <i>is</i> modified <i>after</i> accessing memory		
[Rn], #offset	EA = Rn Rn = Rn + offset	LDR R0, [R1], #4
[Rn], Rm	EA = Rn Rn = Rn + Rm	LDR R0, [R1], R2
[Rn], Rm, LSL #shift	EA = Rn Rn = Rn + (Rm << shift)	LDR R0, [R1], R2, LSL #2
[Rn], Rm, LSR #shift	EA = Rn Rn = Rn + (Rm >> shift)	LDR R0, [R1], R2, LSR #2

* EA is the *Effective Address*, which is the memory address to which the load or store operation is applied.

Note: only a subset of the above addressing modes can be used to load or store halfwords, signed-halfwords or signed-bytes. See ARM Architecture Reference Manual section A5.3.

STM and LDM – STore and Load Multiple

Instruction	Example
Base register Rn is <i>not</i> modified (no !)	
STMmode Rn, {list}	STMIA R12, {R0-R3} Store the contents of R0-R3 in memory at the address contained in R12
LDMmode Rn, {list}	LDMIA R12, {R5,R7,R10} Load R5, R7 and R10 with the contents of memory at the address contained in R12
Base register Rn is modified (Rn!)	
STMmode Rn!, {list}	STMFD SP!, {R4-R12,R14} Push R4-R12 and R14 on to the system stack, updating the system stack pointer
LDMmode Rn!, {list}	LDMFD SP!, {R4-R12,R14} Pop 10 words off the top of the system stack into R4-R12 and R14, updating the system stack pointer

mode

STM – STore Multiple		LDM – Load Multiple	
Instruction	Stack-Oriented Synonym	Instruction	Stack-Oriented Synonym
STMDB (decrement before)	STMFD / PUSH (full descending)	LDMIA (increment after)	LDMFD / POP (full descending)
STMIB (increment before)	STMFA (full ascending)	LDMDA (decrement after)	LDMFA (full ascending)
STMDA (decrement after)	STMED (empty descending)	LDMIB (increment before)	LDMED (empty descending)
STMIA (increment after)	STMEA (empty ascending)	LDMDB (decrement before)	LDMEA (empty ascending)

Example: Sum (using Register Offset Addressing)

23

Design and write an assembly language program that will calculate the sum of 10 word-size values stored in memory

```
LDR    R1, =TESTDATA ; address = address of first word-wise value
LDR    R0, =0          ; sum = 0
LDR    R4, =0          ; count = 0
LDR    R5, =0          ; offset = 0

L1
  CMP    R4, #10        ; while (count < 10)
  BHS    L2             ; {
  LDR    R6, [R1, R5]    ; num = Memory.Word[address + offset]
  ADD    R0, R0, R6      ; sum = sum + num
  ADD    R5, R5, #4      ; offset = offset + 4
  ADD    R4, R4, #1      ; count = count + 1
  B      L1             ; }
L2      ;
```

Scaled Register Offset Addressing

24

Addressing Mode Syntax	Operation	Example
[Rn, Rm, LSL #shift]	$EA = Rn + (Rm \times 2^{\text{shift}})$	LDR R0, [R1, R2, LSL #2]

Effective Address is calculated by adding offset in **Rm**, shifted left by **shift** bits, to the address in the base register **Rn**

Base register **Rn** and offset register **Rm** are not changed

Example: load three consecutive word values from memory into registers R1, R2 and R3 beginning at the address in R0

```
LDR    r0, =label           ; Initialise base register
LDR    r4, =0                ; Initialise index register = 0
LDR    r1, [r0, r4, LSL #2]  ; r1 = Memory.Word[r0 + r4 * 4]
ADD    r4, r4, #1            ; r4 = r4 + 1
LDR    r2, [r0, r4, LSL #2]  ; r2 = Memory.Word[r0 + r4 * 4]
ADD    r4, r4, #1            ; r4 = r4 + 1
LDR    r3, [r0, r4, LSL #2]  ; r3 = Memory.Word[r0 + r4 * 4]
```


Example: Sum (using Scaled Register Offset Addressing)

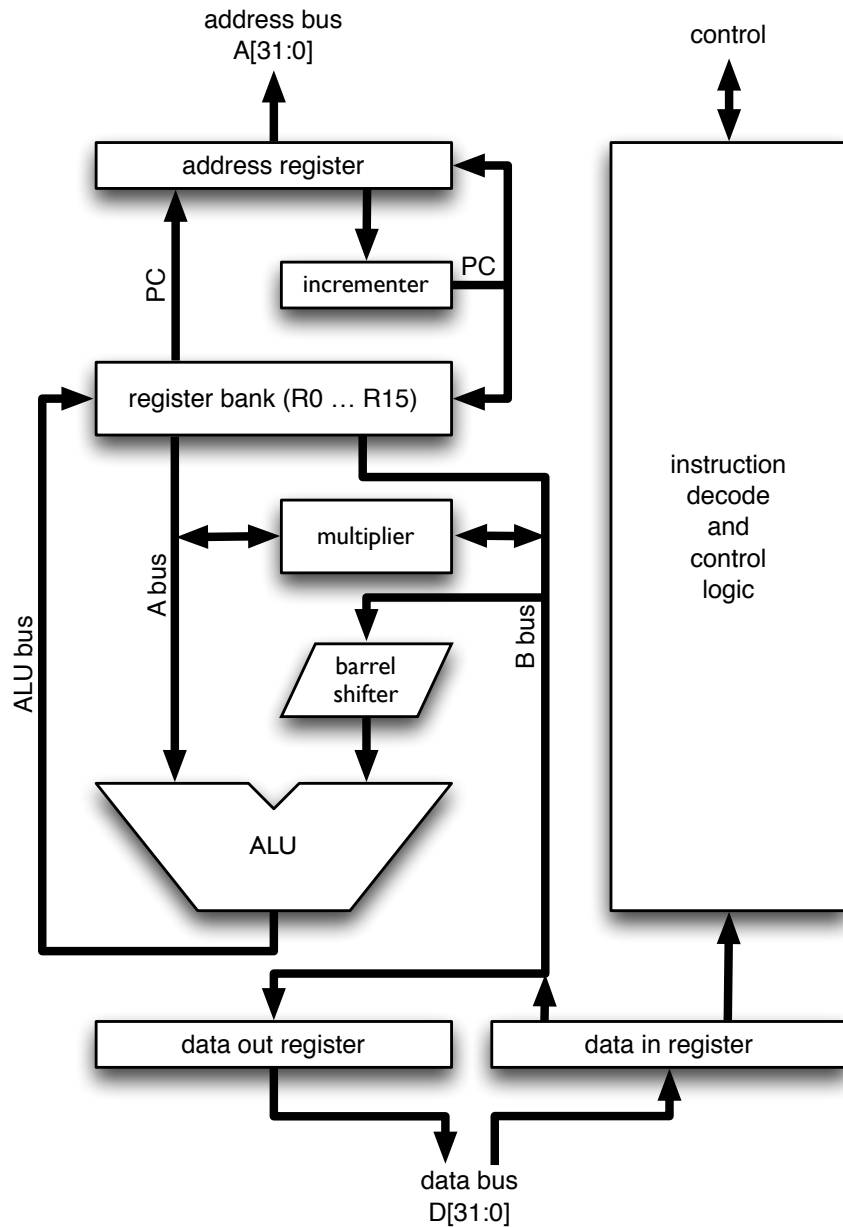
25

Re-write our program to calculate the sum of 10 word-size values stored in memory, this time using scaled register offset addressing

Allows *count* to be used for *offset*!!

```
LDR    R1, =TESTDATA        ; address = address of first word-wise value
LDR    R0, =0                ; sum = 0
LDR    R4, =0                ; count = 0

L1
  CMP   R4, #10              ; while (count < 10)
  BHS   L2                   ; {
  LDR   R6, [R1, R4, LSL #2] ; num = Memory.Word[address + count * 4]
  ADD   R0, R0, R6           ; sum = sum + num
  ADD   R4, R4, #1           ; count = count + 1
  B     L1                   ; }
L2                                           ;
```



Steve Furber, ***“ARM System-on-Chip Architecture”***, 2nd edition, Addison Wesley Professional, 2000.

Pre- and Post-Indexed Addressing

27

Many programs iterate sequentially through memory (examples?)

Often manifested as an LDR/STR, followed by an ADD


```
LDR    R4, [R1]    ; load
ADD     R1, R1, #4   ; increment base address register R1
```

ARM architecture provides a set of addressing modes that incorporate the increment/decrement into the execution of the LDR/STR instruction

Pre-Indexed Addressing	Post-Indexed Addressing
1. Increment / Decrement base address register (Rn)	1. Compute Effective Address
2. Compute Effective Address	2. Increment / Decrement base address register (Rn)

Post-Indexed Addressing Modes

28




Syntax: post-indexed addressing modes place the immediate/register/scale addressing mode operands **after** the []

Behaviour: (i) the LDR/STR is performed first using the original base register value, (ii) the base register is updated by applying the post-index operation

Modes: Immediate Post-Indexed, Register Post-Indexed, Scaled Register Post-Indexed (LSR/LSL)

Pre-Indexed Addressing Modes

29



Syntax: pre-indexed addressing modes place the immediate/register/scale addressing mode operands **inside** the []

Behaviour: (i) the base register is updated by applying the pre-index operation, (ii) the LDR/STR is performed using the updated base register value

Modes: Immediate Pre-Indexed, Register Post-Indexed, Scaled Register Post-Indexed (LSR/LSL)

Example: Upper Case String (using Immediate Post-Indexed Addressing)

30

```
LDR    r1, =TESTSTR        ; address = teststr

L1
LDRB   r0, [r1], #1         ; while ( (char = Memory.Byte[address++])
CMP     r0, #0              ;     != 0 )
BEQ     L3                  ; {
CMP     r0, #'a'            ; if (char >= 'a'
BCC     L2                  ; AND
CMP     r0, #'z'            ; char <= 'z')
BHI     L2                  ; {
BIC     r0, #0x00000020      ; char = char AND NOT 0x00000020
STRB    r0, [r1, #-1]        ; Memory.Byte[address - 1] = char
L2
B       L1                  ;
L3
; }
```

Immediate Post-Indexed Addressing

Note the use of Immediate Offset Addressing to (temporarily) compensate for the earlier post-increment

Example: Sum (using Immediate Post-Indexed Addressing)

31

```
LDR    R1, =TESTDATA ; address = address of first word-wise value
LDR    R0, =0          ; sum = 0
LDR    R4, =0          ; count = 0

L1
  CMP    R4, #10        ; while (count < 10)
  BHS    L2              ; {
  LDR    R6, [R1], #4    ; num = Memory.Word[address]; address=address+4
  ADD    R0, R0, R6      ; sum = sum + num
  ADD    R4, R4, #1      ; count = count + 1
  B      L1              ; }
L2
  ;
```

Review of Single-Dimensional Arrays

Arrays

33

Nothing new here ... you have already been using arrays!

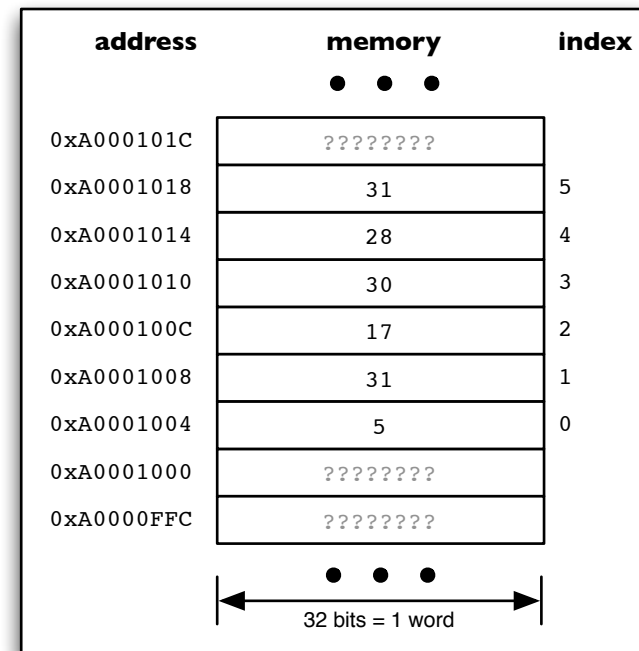
Array – an ordered collection of elements stored sequentially in memory

e.g. integers, ASCII characters, lottery numbers

Homogeneous elements?

(at least with respect to size)

Dimension: number of elements in array



Accessing (loading/storing) Array Elements

34

Step 1: translate array index into byte offset from start address of array in memory

$$\text{<byte offset>} = \text{<index>} \times \text{<elem size>}$$

Step 2: add byte offset to array base address to access element

$$\text{<address>} = \text{<array start address>} + \text{<byte offset>}$$

Example: retrieve the 4th element (index=3) of an array of words

Efficient implementation of random access using **Scaled Register Offset** addressing mode:

```
LDR    r4, =TESTARR          ; pArr = start address of array
LDR    r5, =3                 ; index = 3 (4th element)
LDR    r6, [r4, r5, LSL #2] ; elem = Memory.Word[pArr + (index * 4)]
```

← array access

Array Access Example - Sum

35

Déjà Vu!

```
LDR    R1, =TESTARR      ; start address of myArray
LDR    R0, =0             ; sum = 0
LDR    R4, =0             ; count = 0

L1
  CMP   R4, #10           ; while (count < 10)
  BHS   L2                ; {
  LDR   R6, [R1, R4, LSL #2] ; num = myArray[count]
  ADD   R0, R0, R6        ; sum = sum + num
  ADD   R4, R4, #1        ; count = count + 1
  B     L1                ; }
L2      ;
```

The pseudo-code
comments have changed but
the program is identical!!!

(See **Addressing Modes**)

Multi-Dimensional Arrays


2D Arrays

37

Arrays can have more than one dimension

e.g. a two-dimensional array – analogous to a table containing elements arranged in rows and columns

Stored in memory by mapping the 2D array into 1D memory, e.g.




2D Arrays

38

Row-major order: 2D array is stored in memory by storing each **row** contiguously in memory

Column-major order: 2D array is stored in memory by storing each **column** contiguously in memory



Rare – row-major order is the assumed norm

2D Array Example

39

2D array declared in memory

```
AREA TestData, DATA, READWRITE

array DCD 6, 3, 8, 2, 5, 2, 9, 1 ; row 0
      DCD 3, 7, 2, 8, 5, 7, 2, 7 ; row 1
      DCD 2, 4, 7, 4, 2, 6, 7, 4 ; row 2
      DCD 1, 9, 3, 2, 9, 5, 6, 8 ; row 3
      DCD 7, 5, 3, 7, 5, 8, 2, 1 ; row 4
      DCD 6, 4, 8, 9, 0, 3, 2, 5 ; row 5
```

... or equivalently ...

```
AREA TestData, DATA, READWRITE

array DCD 6, 3, 8, 2, 5, 2, 9, 1, 3, 7, 2, 8, 5, 7, 2, 7, 2, 4, 7, 4, 2
      DCD 6, 7, 4, 1, 9, 3, 2, 9, 5, 6, 8, 7, 5, 3, 7, 5, 8, 2, 1, 6, 4
      DCD 8, 9, 0, 3, 2, 5
```

Why are these
equivalent?

Accessing 2D Array Elements

40

Example: retrieve the element at the 3rd row and 2nd column of a 2D array of words with 6 rows and 8 columns – **array[3][2]**

Step 1: translate 2D array index into 1D array index

$$\text{<index>} = (\text{<row>} \times \text{<row size>}) + \text{<col>}$$

Step 2: translate 1D array index into byte offset from start address of array in memory

$$\text{<byte offset>} = \text{<index>} \times \text{<elem size>}$$

Step 3: add byte offset to array base address to access element

$$\text{<address>} = \text{<array base address>} + \text{<byte offset>}$$

Step 1 is new

Steps 2 & 3 are the same as those for a 1-D array ...

... because our 2-D array is just a different interpretation of a 1-D array!

2D Array Access Example

41

Example: retrieve the element at the 3rd row and 2nd column of a 2D array of words with 6 rows and 8 columns – array[3][2]

```
LDR    r4, =TESTARR          ; pArr = address of array start
LDR    r5, =COLSZ             ; load col_size
LDR    r6, =ROWSZ             ; load row_size

; looking for array[3][2] (4th row, 3rd column)

LDR    r1, =3                 ; row = 3
LDR    r2, =2                 ; col = 4

; <byte offset> = ((row * <row_size>) + col) * <elem size>

MUL    r7, r1, r6             ; index = row * row_size
ADD    r7, r7, r2             ; index = index + col
LDR    r0, [r4, r7, LSL #2] ; elem = Memory.Word[ pArr + (index*4) ]
```

index
calculation

array
access

Multi-dimension arrays

42

e.g. a 3D array of size $sz \times sy \times sx$

In general, the index of element $a[z][y][x]$ is:

$$\text{index} = ((z \times sy \times sx) + (y \times sx) + x)$$

e.g. a 4D array of size $sz \times sy \times sx \times sw$

In general, the index of element $a[z][y][x][w]$ is:

$$\text{index} = ((z \times sy \times sx \times sw) + (y \times sx \times sw) + (x \times sw) + w)$$

Warning: an array of bytes with odd dimensions may cause the data following the array to be odd aligned, requiring padding of one byte (similarly for half-words)